

Teaching an AI to Double-Check Its Own Work

Review Residuals — an illustrated, plain-English companion

Kyle Kramer

NeraTech LLC

June 2026

In one sentence. Today’s AI models accept every internal step blindly; we added a small “check it before you commit to it” step—borrowed from how pilots, surgeons, and nuclear operators avoid mistakes—and found that it does nothing for small models but, as the models grow, starts to win, until at the largest sizes we tested it beats the standard method today’s AI actually uses.

This companion walks through the same idea, figures, and results as our technical paper, but in everyday language. Every figure and table from the research version is here, each with a plain explanation of what it shows and why it matters.

1. What kind of AI are we talking about?

The AI behind tools like ChatGPT is a *language model*. At its core it does one simple thing: read some text and predict what comes next, one piece at a time. To do that well it passes the text through many stacked *layers*, each refining the model’s internal sense of things before handing it on. Picture a long assembly line: a sentence goes in one end, passes through dozens of stations, and a polished prediction comes out the other.

2. The hidden problem: the AI trusts every step blindly

At every station on that assembly line, the layer *proposes a change* to the work-in-progress—and that change is **applied automatically, in full, every time**. There is no inspection step. Nothing pauses to ask, “is this particular change actually a good idea?” Some changes

are great, some are noise; the machine trusts them all equally. It works remarkably well anyway—but it’s a strange way to build something this important. (The technical name for the wiring that does this is a “residual connection”; the plain version is: every step is accepted, no questions asked.)

3. The fix comes from aviation, surgery, and nuclear power

Those fields solved this exact problem long ago—and the reason they had to is worth sitting with. Even brilliant, fully-attentive experts miss things. In a famous experiment, people asked to count basketball passes completely failed to notice a person in a gorilla suit stroll through the middle of the scene [18]. The lesson isn’t that those people were dim; it’s that *every* mind, human or machine, has hard limits on what it can attend to and get right—intelligence itself is *bounded*, and you can’t fix that by demanding more of it [19]. So high-reliability fields engineer checks *around* the expert, and they even have a name for the toolkit: *human-performance (HP) tools*. Pilots read a step back before doing it (self-checking); a second person confirms critical actions (peer-checking); and, most relevant here, a proposed action is checked before it is committed—*independent verification* [22]. A simple pre-surgery checklist built on these ideas measurably cut deaths. The common thread is never “hire smarter people”—it is *build a check into the process itself*. We asked: what if we built one of those HP tools—independent verification—into the AI?

4. Our idea: a built-in “trust dial”

Before each layer’s proposed change is applied, we add a small learned *trust dial*. It looks at two things—where the work stands now, *and* the specific change being proposed—and decides how much of that change to let through. That second part is what’s new.

How to read Figure 1: in all three pictures the vertical line is the “work so far,” flowing upward untouched. The box on the right is the layer proposing a change. The difference is the dial: the standard approach has none (it adds the whole change); Highway’s dial peeks only at the work so far; ours (the red arrow) is the only one that actually *looks at the proposed change itself* before deciding how much to trust it—like an inspector who picks up the part and examines it, rather than just glancing at the line.

All three keep the identity path intact (additive). Review’s gate is the only one that inspects the proposed update u_l .




Figure 1: (Figure 1 from the paper.) Three ways to handle a proposed change, all keeping the original work intact. (a) Standard: accept the change fully, always. (b) Highway, the closest existing method: a dial that looks only at the current situation. (c) Review Residuals (ours): a dial that looks at both the situation *and* the specific proposed change (the red arrow).

5. How our idea compares to the others

Researchers have tried several ways to control how much of each step to apply. Table 1 lays them out. Only the last column matters for what’s new here.

6. A wall we hit—and climbed

This part shows the work was done carefully. Our first version of the dial used a particular bit of math that worked fine for small models but, at large sizes, **would not train at all**—the model got stuck and learned almost nothing. We diagnosed exactly why (the math was starving the deepest layers of the learning signal they need—the same problem a famous 2015 breakthrough was invented to fix [2]) and corrected it with a principled change that keeps the original work flowing through untouched. After the fix, the large models trained cleanly. Finding and fixing that wall is itself a useful result.

Method		What it adjusts	Reacts to the input?	Looks at the proposed change?
Standard (what AI uses today)		nothing — adds it all	no	no
ReZero / Fixup / Layer-Scale		a fixed “volume knob”	no	no
Highway [3]		a dial set by the situation	yes	no
Attention Residuals [8]		blends across many steps	yes	not directly
Review (ours)	Residuals	a dial set by situation <i>and</i> change	yes	yes

Table 1: (Table 1 from the paper, in plain terms.) Every prior method either ignores the input or reacts only to the current situation. Ours is the only one that looks at the *specific change being proposed* before deciding how much of it to apply.

7. How we tested it (a fair race)

An AI model starts knowing nothing; we show it a big pile of text and have it practice predicting the next piece, adjusting itself millions of times. To score it, we test on text it has never seen and measure how *surprised* it is—less surprise is better, so throughout, **lower numbers are better**. We built the competing methods too and trained them all under identical conditions, ran each several times from different starting points to rule out luck, and—crucially—gave the competitors *extra size* so they had at least as many “parts” as ours. That last point matters: it means any win can’t just be because our method is bigger.

8. What we found

We trained five sizes, from small up to one billion parameters. Table 2 has the actual scores, and Figure 2 shows the pattern.

How to read Figure 2: the horizontal axis is model size (small on the left, a billion on the right). The vertical axis is how much better our method is—above the dashed line we win, below it the other method wins. The grey line (versus the plain standard method) literally *starts below the line*—the standard beats us on small models—then climbs and ends well above it at a billion parameters. The purple line (versus the alternative dial) does the same. In plain terms: **the benefit emerges as the models get bigger**. That’s unusual and encouraging—most clever tweaks get *harder* to justify as models grow; this one got *easier*.

Model size	Ours (Review)	Highway	Standard	Did ours win?
60M (small)	1.6891	1.6923	1.6805	no — standard was better
150M	1.5576	1.5625	1.5548	tie
320M	1.5001	1.5042	1.5037	tie (barely ahead)
590M	1.4795	1.4889	1.4901	yes, clearly (statistically solid)
1B (largest)	1.4876	1.5031	1.5040	yes, by more (strong trend)

Table 2: (Table 2 from the paper, in plain terms.) Scores are “surprise” — lower is better. Read top to bottom and watch it change: at the smallest size the plain standard method actually beats ours; through the middle it’s a tie; then at the two largest sizes ours pulls clearly ahead of *both* the alternative and the standard method.

9. What’s encouraging — and what we’re not claiming

What’s encouraging is real. This is a fresh idea, drawn from a safety tradition almost nobody has applied to AI. At the largest scale we could test, it beats the *standard* method—not just an alternative—and the advantage grows with size.

And what to keep in perspective. These are still small models by industry standards, the winning margin is about one percent, the billion-parameter result is a strong trend that points the same way as the clear win just below it, and we tested on one simple kind of text [11]. None of that is a strike against the idea—it is simply where a promising early result sits. What makes it worth pursuing is the direction: as the models grew, the advantage did not fade, it *grew*.

10. Why it could matter

As AI takes on more, two things become more valuable: models that are *reliable* and models we can *understand*. This nudges both—it gives a model a habit of checking its own work, and the dial’s settings are something you can actually inspect. There’s a century of hard-won knowledge about making high-stakes systems trustworthy, from cockpits to operating rooms, and almost none of it has reached how we build AI. This is one small, concrete attempt to bring it in—and, encouragingly, the bigger the model, the more it seemed to help.

Questions you might have

Does this make the AI self-aware or able to improve itself? No. It’s a quality-control step inside the existing process; the model checks each step before applying it. It does not redesign itself.




Figure 2: (Figure 2 from the paper.) How far ahead our method is, as the model grows (left to right). Above the dashed line = ours is better; below = the other method is better. Both lines start at or below the line for small models and rise clearly at the largest sizes. The star marks where the win is statistically solid.

Is it a brand-new invention? The idea of “dials” on these steps has been explored before, but the key move here is new: having the dial look at the *specific change* being proposed, not just the situation. As far as we know, no prior method does that—and building safety-engineering verification into an AI’s architecture is new as well.

Will my chatbot get better tomorrow? Not directly yet—but here is the encouraging part: the benefit grew with *every* size we tested. That is exactly the pattern you would want if it is going to matter for the large models behind real products. Whether it keeps growing all the way up is the next experiment, and it is a very answerable one.

The takeaway. We borrowed a safety idea—check before you commit—and built it into an AI. As the models grew, it went from no help to clearly winning: at the largest sizes we tested, it beat *both* the alternative and the standard method today’s AI uses—and the lead *widened* with scale. That upward trend is the exciting part: it points straight at the large models that matter in practice, and following it there is the obvious next step.

References

- [1] A. Vaswani et al. Attention Is All You Need. *NeurIPS*, 2017. arXiv:1706.03762.

- [2] K. He, X. Zhang, S. Ren, J. Sun. Deep Residual Learning for Image Recognition. *CVPR*, 2016. arXiv:1512.03385.
- [3] R. K. Srivastava, K. Greff, J. Schmidhuber. Highway Networks. arXiv:1505.00387, 2015.
- [4] T. Bachlechner et al. ReZero is All You Need. *UAI*, 2021. arXiv:2003.04887.
- [5] H. Touvron et al. Going Deeper with Image Transformers (CaiT / LayerScale). *ICCV*, 2021. arXiv:2103.17239.
- [6] H. Zhang, Y. N. Dauphin, T. Ma. Fixup Initialization. *ICLR*, 2019. arXiv:1901.09321.
- [7] H. Wang et al. DeepNet: Scaling Transformers to 1,000 Layers. arXiv:2203.00555, 2022.
- [8] Kimi Team. Attention Residuals. arXiv:2603.15031, 2026.
- [9] B. Zhang, R. Sennrich. Root Mean Square Layer Normalization. *NeurIPS*, 2019. arXiv:1910.07467.
- [10] I. Loshchilov, F. Hutter. Decoupled Weight Decay Regularization (AdamW). *ICLR*, 2019. arXiv:1711.05101.
- [11] R. Eldan, Y. Li. TinyStories. arXiv:2305.07759, 2023.
- [12] A. Graves. Adaptive Computation Time for Recurrent Neural Networks. arXiv:1603.08983, 2016.
- [13] A. Banino, J. Balaguer, C. Blundell. PonderNet: Learning to Ponder. arXiv:2107.05407, 2021.
- [14] N. Shazeer et al. Outrageously Large Neural Networks: The Sparsely-Gated MoE Layer. *ICLR*, 2017. arXiv:1701.06538.
- [15] D. Raposo et al. Mixture-of-Depths. arXiv:2404.02258, 2024.
- [16] R. Xiong et al. On Layer Normalization in the Transformer Architecture. *ICML*, 2020. arXiv:2002.04745.
- [17] J. L. Ba, J. R. Kiros, G. E. Hinton. Layer Normalization. arXiv:1607.06450, 2016.
- [18] D. J. Simons, C. F. Chabris. Gorillas in Our Midst: Sustained Inattentional Blindness for Dynamic Events. *Perception*, 28(9):1059–1074, 1999.
- [19] K. Kramer. *The Operator’s Guide to AI Agents*. NeraTech LLC, 2026. (Bounded-intelligence thesis.)
- [20] J. Reason. *Human Error*. Cambridge University Press, 1990.
- [21] K. E. Weick, K. M. Sutcliffe. *Managing the Unexpected*. Jossey-Bass, 2007.
- [22] U.S. Department of Energy. *Human Performance Improvement Handbook*, DOE-HDBK-1028-2009, 2009.